

-
-
-
-
-
-
-
-
-

Chapter IV

INTERPROCESS COMMUNICATION

Jehan-François Pâris
jfparis@uh.edu

CS@UH

-
-
-
-
-

•
•
•

Chapter overview

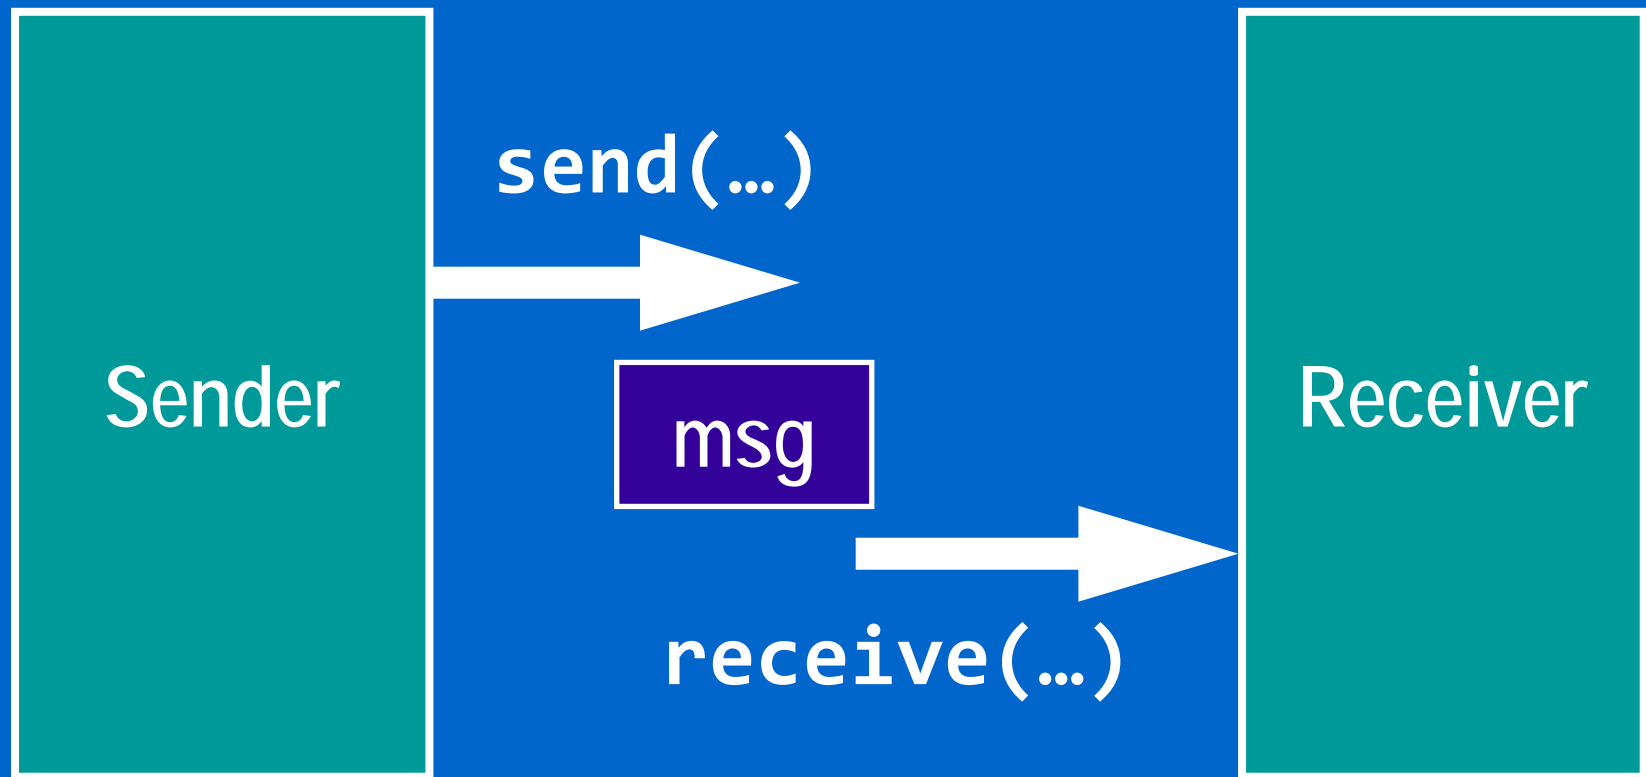
- *Types of IPC*
 - Message passing
 - Shared memory
- *Message passing*
 - Blocking/non-blocking, ...
 - Datagrams, virtual circuits, streams
 - Remote procedure calls

•
•
•

Message passing (I)

- Processes that want to exchange data send and receive *messages*
- Any *message exchange* requires
 - *One send*
`send(addr, msg, length);`
 - *One receive*
`receive(addr, msg, length);`

Message passing (II)



-
-
-

Advantages

- *Very general*
 - Sender and receivers can be on different machines
- *Relatively secure*
 - Receiver can inspect the messages it has received before processing them

-
-
-

Disadvantages

- *Hard to use*
 - Every data transfer requires a **send()** and a **receive()**
 - Receiving process must *expect* the **send()**
 - Might require forking a special thread

•
•
•

Shared Memory

- Name says it
 - Two or more processes share a part of their address space



-
-
-

Advantages

- *Fast and easy to use*
 - The data are there
- but
 - Some concurrent accesses to the shared data can result into small disasters
 - Must synchronize access to shared data
 - Topic will be covered in next chapter

•
•
•

Disadvantages

- *Not a general solution*
 - Sender and receivers must be on the *same machine*
- *Less secure*
 - Processes can directly access a part of the address space of other processes

•
•
•

MESSAGE PASSING

- *Defining issues*
 - Direct/Indirect communication
 - Blocking/Non-blocking primitives
 - Exception handling
 - Quality of service
 - Unreliable/reliable datagrams
 - Virtual circuits, streams

•
•
•


Direct communication (I)

- Send and receive system calls always specify **processes** as destination or source:
 - **send(process, msg, length);**
 - **receive(process, msg, &length);**
- Most basic solution because there is
 - No intermediary between sender and receiver

•
•
•

An analogy

- Phones without switchboard
 - Each phone is hardwired to another phone



•
•
•

Direct communication (II)

- Process executing the receive call must know the identity of all processes likely to send messages
 - Very bad solution for *servers*
 - *Servers have to answer requests from arbitrary processes*

•
•
•


Indirect communication (I)

- Send and receive primitives now specify an *intermediary entity* as destination or source: the *mailbox*
send(mailbox, msg, size);
receive(mailbox, msg, &size);
- Mailbox is a system object created by the kernel at the request of a user process

•
•
•

An analogy (I)

- Phones with a switchboard
 - Each phone can receive calls from any other phone



•
•
•

An analogy (II)

- Each phone has now a *phone number*
 - Callers dial that number, not a person's name
- Taking our phone with us allows us to receive phone calls *from everybody*

•
•
•

Indirect communication (II)

- Different processes can send messages to the same mailbox
 - A process can receive messages from processes it does not know anything about
 - A process can wait for messages coming from different senders
 - Will answer the first message it receives

-
-
-

Mailboxes

- Mailboxes can be
 - *Private*
 - Attached to a specific process
 - *Think of your cell phone*
 - *Public*
 - System objects
 - *Think of a house phone*

-
-
-

Private mailboxes

- Process that requested its creation and its children are the only processes that can receive messages through the mailbox are that process and its children
- Cease to exist when the process that requested its creation (and all its children) terminates.
- Often called **ports**
- *Example:* BSD *sockets*

•
•
•

Public mailboxes

- Owned by the system
- Shared by all the processes having the right to receive messages through it
- Survive the termination of the process that requested their creation
- Work best when all processes are on the same machine
- *Example:* System V UNIX *message queues*

•
•
•

Blocking primitives (I)


- A *blocking send* does not return until the receiving process has received the message
 - No *buffering* is needed
 - *Analogous to what is happening when you call somebody who does not have voice mail*

•
•
•

Blocking primitives (II)

- A *blocking receive* does not return until a message has been received
- *Like waiting by the phone for an important message or staying all day by your mailbox waiting for the mail carrier*

Blocking primitives (III)



•
•
•

Non-blocking primitives (I)

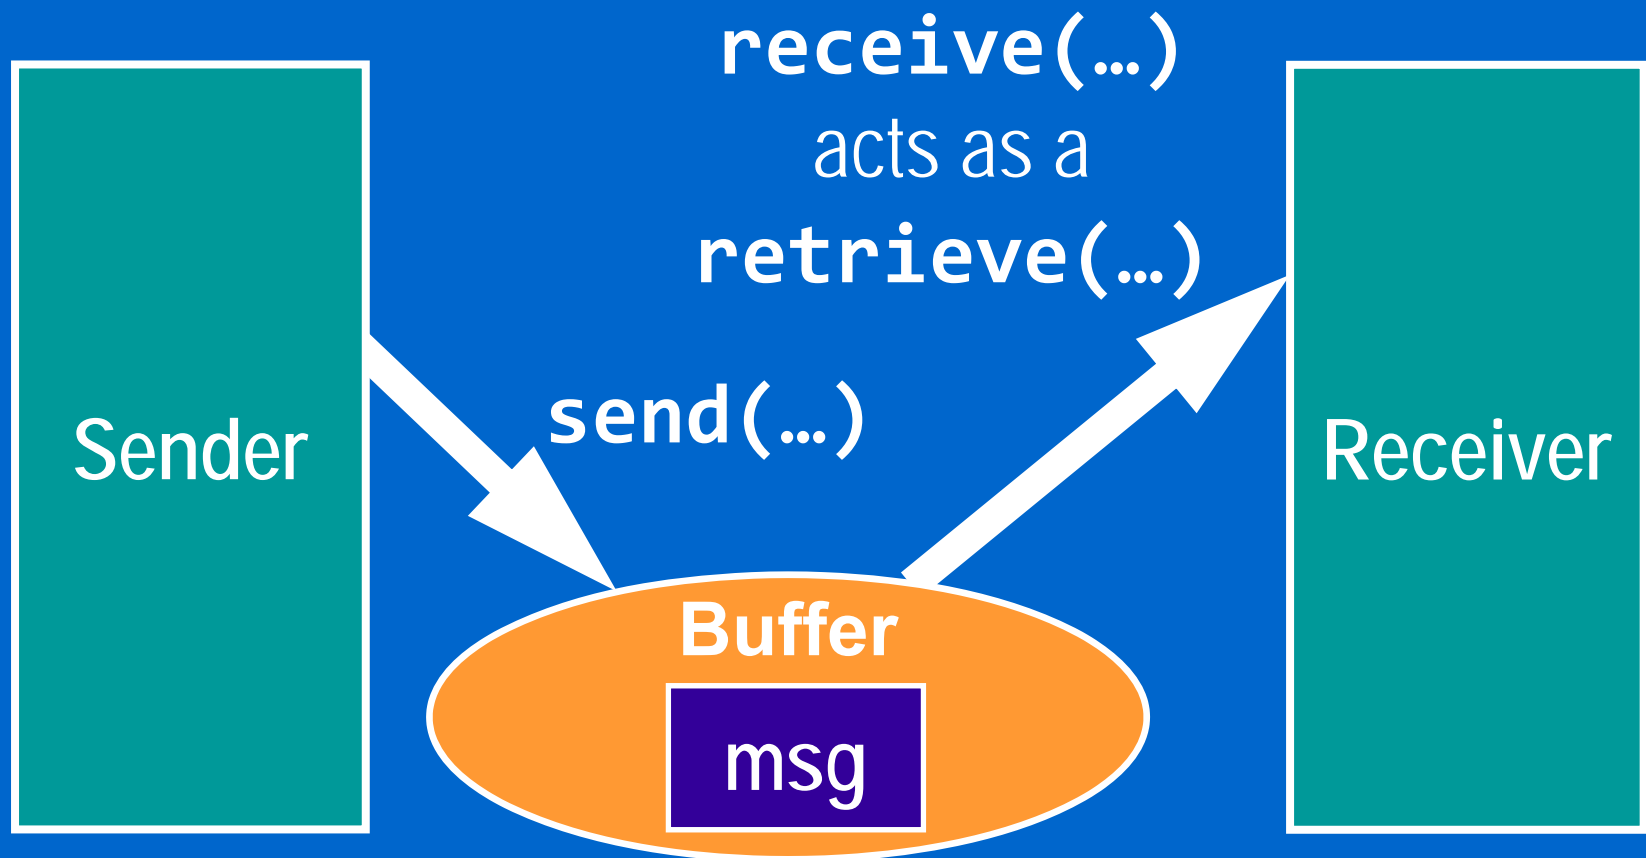
- A *non-blocking send* returns as soon as the message has been accepted for delivery by the OS
 - Assumes that the OS can store the message in a *buffer*.
 - *Like mailing a letter: once the letter is dropped in the mailbox, we are **done**.* The mailbox will hold your letter until a postal employee picks it up

•
•
•

Non-blocking primitives (II)

- A *non-blocking receive* returns as soon as it has either retrieved a message or learned that the mailbox is empty
 - *Like checking whether your mail has arrived or not*

Non-blocking primitives (III)



•
•
•

Simulating blocking receives

- Can simulate a blocking receive with a non-blocking receive within a loop:

```
do {  
    code = receive(mbox,msg,size);  
    sleep(1); // delay  
} while (code == EMPTY_MBOX);
```

– Known as a *busy wait*

•
•
•

Simulating blocking sends

- Can simulate a blocking send with two non-blocking sends and a blocking receive:-
 - Sender sends message and requests an acknowledgement (ACK)
 - Sender waits for ACK from receiver using a blocking receive
 - Receiver sends ACK

-
-
-

The standard choice

- In general we prefer
 - *Indirect naming*
 - *Non-blocking sends*
 - Sender does not care about what happens once the message is sent
 - *Similar to UNIX delayed writes*
 - *Blocking receives*
 - Receiver needs the data to continue

•
•
•

Buffering

- *Non-blocking primitives* require *buffering* to let OS store somewhere messages that have been sent but not yet received
- These buffers can have
 - *Bounded capacity*
 - Refuse to receive messages when the buffer is full
 - Theoretically *unlimited capacity*.

•
•
•

An explosive combination (I)

- *Blocking receive* does not go well with *direct communication*
 - Processes cannot wait for messages from several sources without using special parallel programming constructs:
 - *Dijkstra's alternative command*

An explosive combination (II)

- Using blocking receives with direct naming does not allow the receiving process to receive any messages from any process but the one it has specified



•
•
•

Exception condition handling

- Must specify what to do if one of the two processes dies
 - Especially important whenever the two processes are on two different machines
 - Must handle
 - *Host failures*
 - *Network partitions*

•
•
•

Quality of service

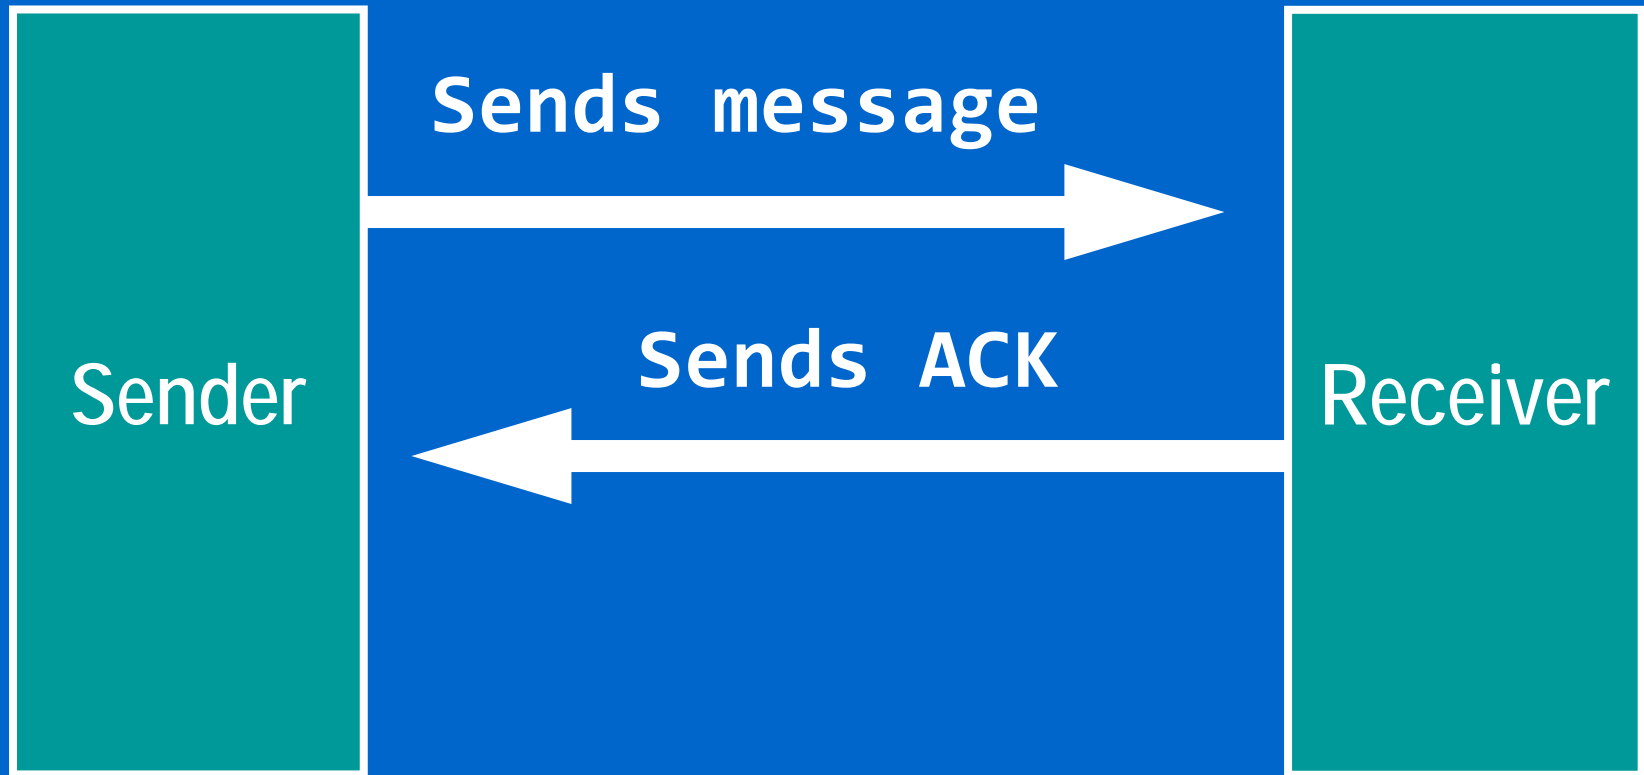
- When sender and receiver are on different machines, messages
 - Can be *lost*, *corrupted* or *duplicated*
 - Arrive *out of sequence*
- Can still decide to provide *reliable message delivery*

•
•
•

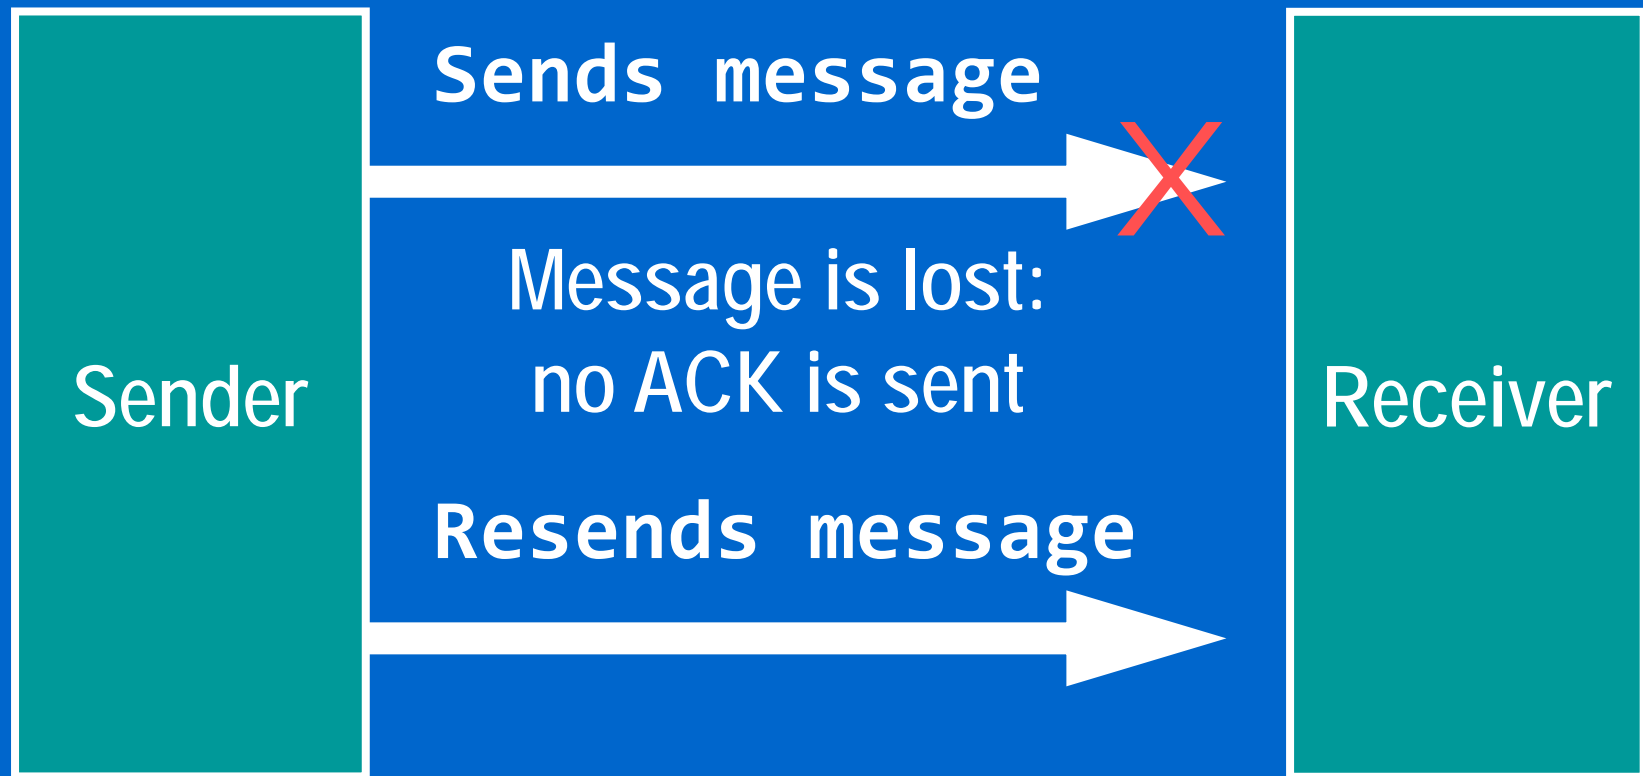
Positive Acknowledgments

- Basic technique for providing reliable delivery of messages
- Destination process sends an *acknowledgment message* (ACK) for every message that was correctly delivered
 - Damaged messages are ignored
- Sender resends any message that has not been acknowledged within a fixed time frame

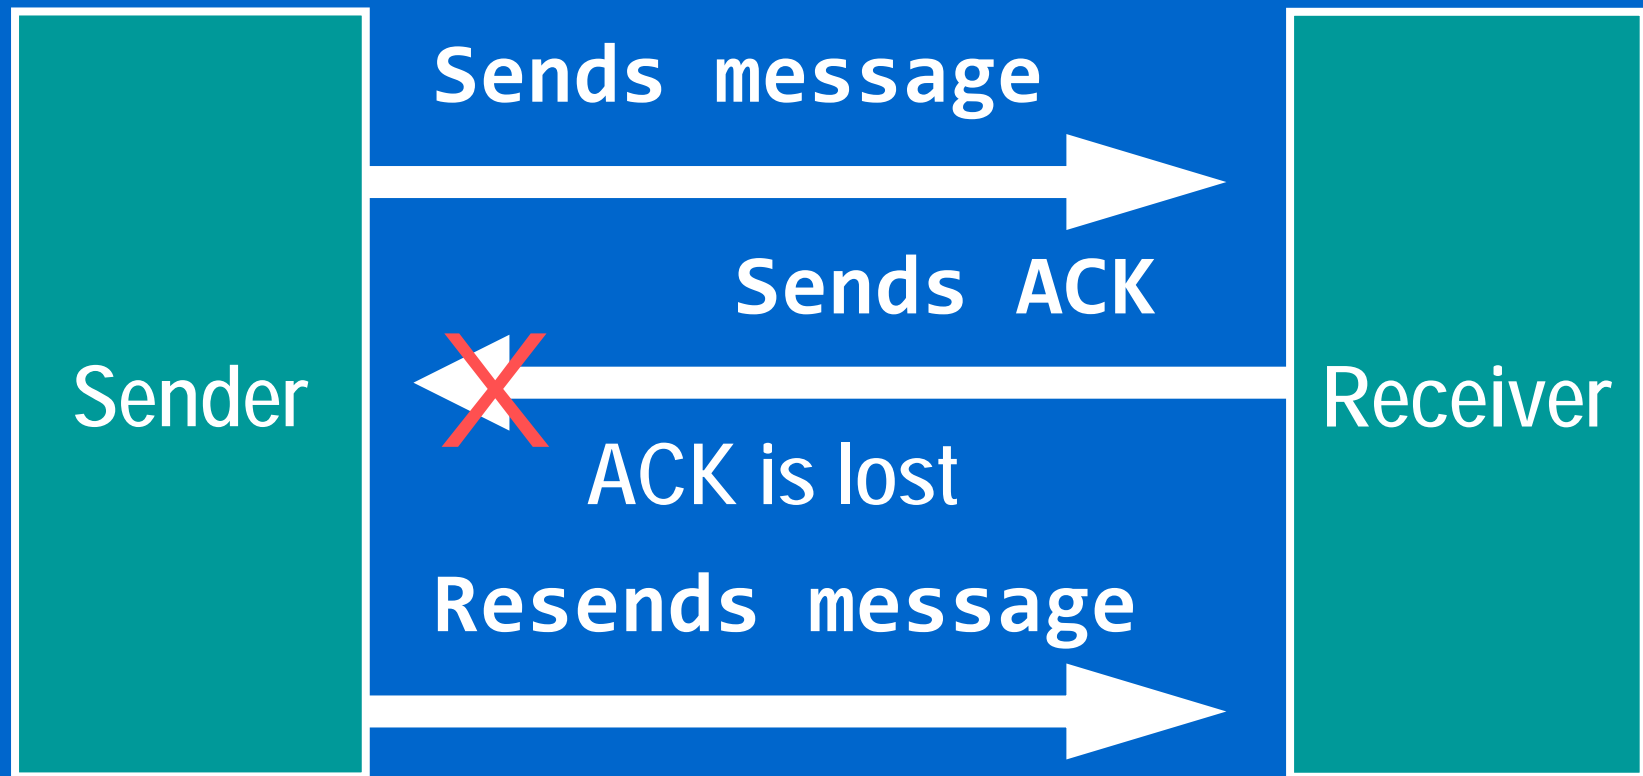
First Scenario



Second Scenario



Third Scenario (I)



-
-
-

Third Scenario (II)

- Receiver *must* acknowledge a second time the message
 - Otherwise it would be resent one more time
- Rule is
 - *Acknowledge any message that does not need to be resent!*

•
•
•

Classes of service

- *Datagrams:*
 - Messages are send one at time
- *Virtual circuits:*
 - Ordered sequence of messages
 - *Connection-oriented* service
- *Streams:*
 - Ordered sequence of bytes
 - Message boundaries are ignored

•
•
•

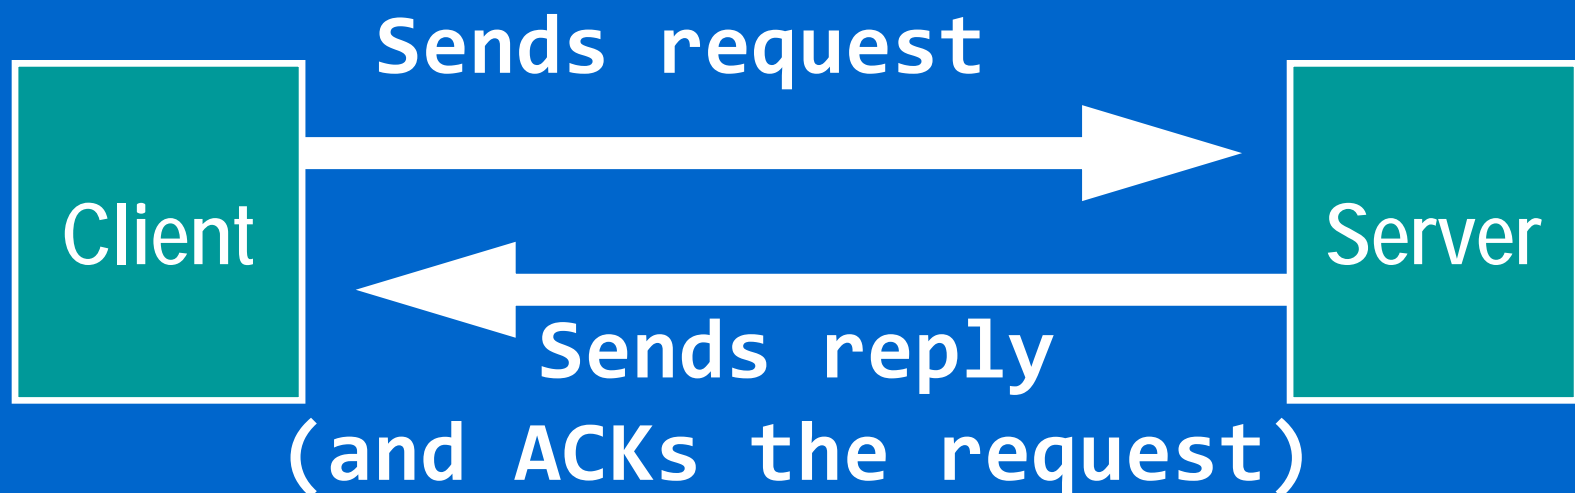
Datagrams

- Each message is sent *individually*
 - Some messages can be *lost*, other *duplicated* or arrive *out of sequence*
 - *Equivalent of a conventional letter*
- *Reliable datagrams:*
resent until they are acknowledged
- *Unreliable datagrams*

•
•
•

Unreliable datagrams (I)

- Messages are not acknowledged
- Works well when message requests a reply
 - Reply is *implicit ACK* of message



•
•
•

Unreliable datagrams (II)

- Exactly what we do in real life:
 - *We rarely ACK emails and other messages*
 - *We reply to them!*
- Sole reason to ACK a request is when it might take a long time to reply to it

-
-
-

UDP

- *User Datagram Protocol*
- Best known datagram protocol
- Provides an unreliable datagram service
 - Messages can be *lost*, *duplicated* or arrive *out of sequence*
- Best for short interactions
 - One request and one reply

•
•
•

Virtual circuits (I)

- Establish a *logical connection* between the sender and the receiver
- Messages are *guaranteed* to arrive in sequence without lost messages or duplicated messages
 - *Analogous to the words of a phone conversation*

-
-
-

Virtual circuits (II)

- Require setting up a virtual connection *before* sending any data
 - Costlier than datagrams
- Best for transmitting large amounts of data that require sending several messages
 - *File transfer protocol* (FTP)
 - *Hypertext transfer protocol* (HTTP)

•
•
•

Streams

- Like virtual circuits
- Do *not* preserve message boundaries:
 - Receiver sees a *seamless stream of bytes*
- Offspring of UNIX philosophy
 - Record boundaries do not count
 - Message boundaries should not count

-
-
-

TCP

- Transmission Control Protocol
- Best known stream protocol
- Provides a reliable stream service
- Said to be *heavyweight*
 - Requires three messages (*packets*) to establish a virtual connection

•
•
•

Datagrams and Streams

- *Datagrams:*
 - Unreliable
 - Not ordered
 - Lightweight
 - Deliver messages
- *Example:*
 - UDP

- *Streams:*
 - Reliable
 - Ordered
 - Heavyweight
 - Stream-oriented
- *Example:*
 - TCP

•
•
•

REMOTE PROCEDURE CALLS

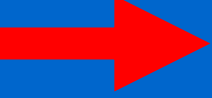
- Apply to **client-server** model of computation
- A typical client-server interaction:

```
send_req(args);      rcv_req(&args);  
                     process(args, &results);  
                     send_reply(results);  
rcv_reply(&results);
```

•
•
•

RPC (cont'd)

- Very similar to a procedure call to a procedure:

- `xyz(args, &results);`  `xyz(...)` {

 return;
} // xyz
 ...

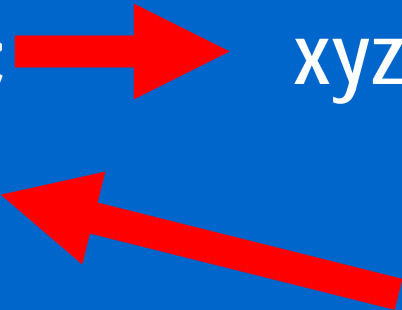
- Try to use the same formalism

•
•
•

RPC (cont'd)

- We could write

```
rpc(xyz, args, &results); → xyz(...) {  
                               ....  
                               return;  
                               } // xyz  
...
```



and let system take care of all message passing details

-
-
-

Advantages

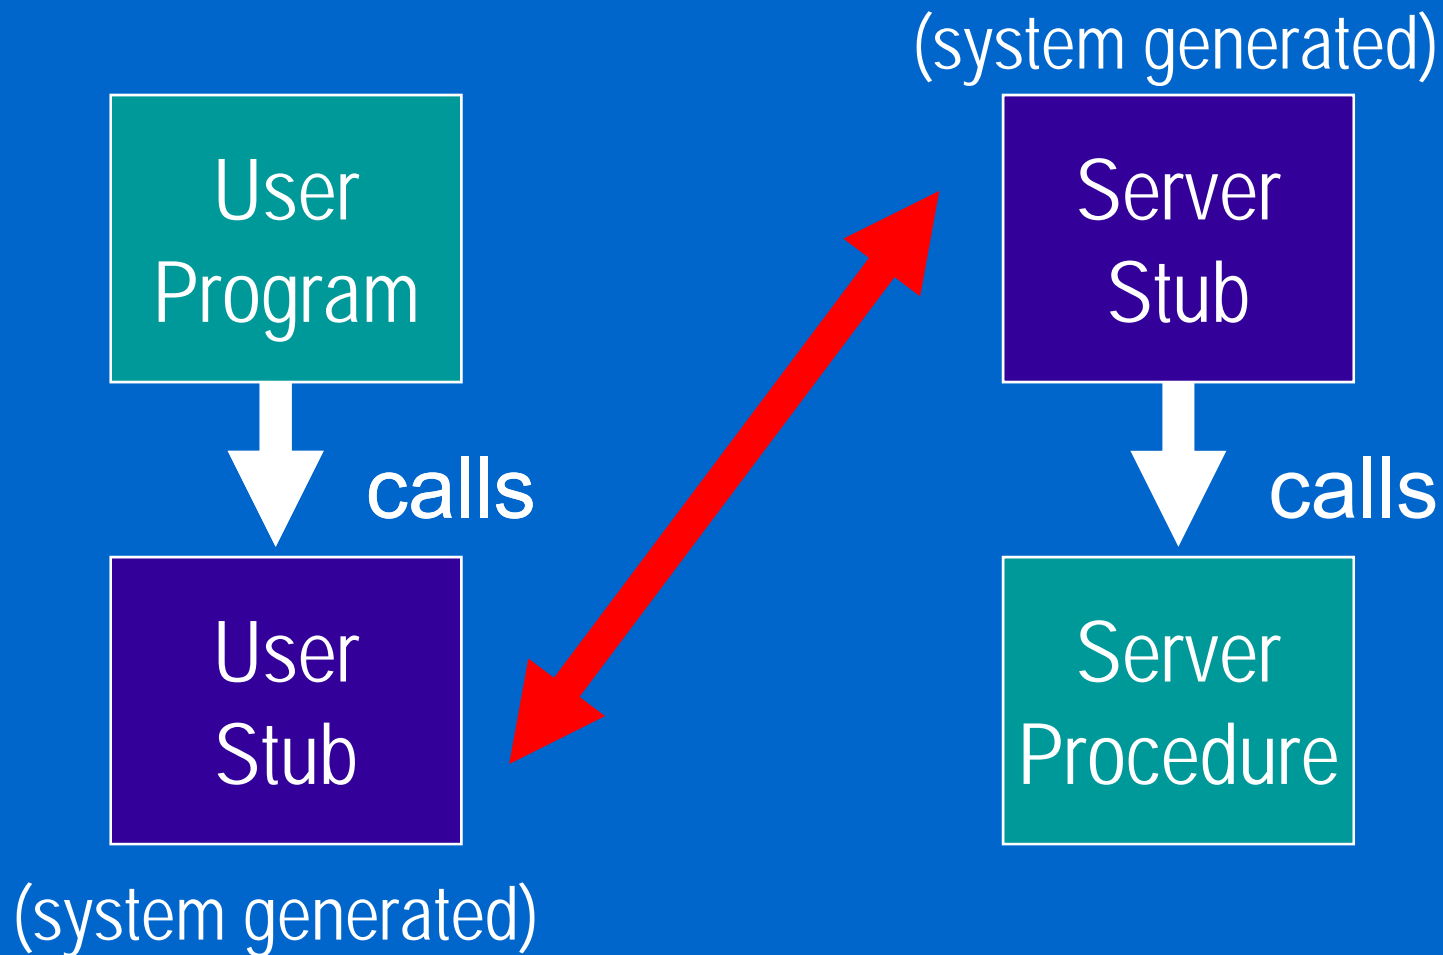
- Hides all details of message passing
 - Programmer can focus on the logic of her application
- Provides a higher level of abstraction
- Extends a well-known model of programming
 - Anybody that can use procedures and function can quickly learn to use remote procedure calls

-
-
-

Disadvantage


- The illusion is not perfect
 - RPCs do not always behave exactly like regular procedure calls
 - Client and server do not share the same address space
- Programmer must remain aware of these subtle and not so subtle differences

General Organization



•
•
•

What the programmer sees



-
-
-

The user program

- Contains the user code
- Calls the user stub

`rpc(xyz, args, &results);`

- *Appears* to call the server procedure

•
•
•

The user stub

- Procedure generated by RPC package:
 - Packs arguments into request message and performs required data conversions (*argument marshaling*)
 - Sends request message
 - Waits for server's reply message
 - Unpacks results and performs required data conversions (*argument unmarshaling*)

•
•
•

The server stub

- Generic server generated by RPC package:
 - Waits for client requests
 - Unpacks request arguments and performs required data conversions
 - Calls appropriate server procedure
 - Packs results into reply message and performs required data conversions
 - Sends reply message

•
•
•

The server procedure

- Procedure called by the server stub
- Written by the user
- Does the actual processing of user requests

•
•
•

Differences with regular PC

- Client and server processes *do not share* the *same address space*
- Client and server can be on *different machines*
- Must handle *partial failures*

-
-
-

No shared address space

- This means
 - *No global variables*
 - *Cannot pass addresses*
 - Cannot pass arguments by reference
 - Cannot pass dynamic data structures through pointers

•
•
•

The solution

- RPC *can pass arguments by value and result*
 - Pass the *current value* of the argument to the remote procedure
 - *Copy* the *returned value* in the user program
- Not the same as passing arguments by reference

Passing by reference

Caller:

```
...  
i = 0;  
abc(&i);  
...
```

i

Procedure abc() will
directly increment
variable i

```
abc(int *k){  
    (*k)++;  
}
```


Passing by value and result

Caller:

```
...  
i = 0;  
abc(&i);  
...
```

i

i = 0

i = 1

Variable i is increased
after caller receives
server's reply

```
abc(int *k){  
    (*k)++;  
}
```

•
•
•

An example (I)

- Procedure doubleincrement

```
doubleincrement(int *p,int *q) {  
    (*p)++ ; (*q)++ ;  
} // doubleincrement
```

- Calling

```
doubleincrement(&m, &m);
```

should increment *m twice*

•
•
•

An example (II)

- Calling

```
doubleincrement(&m, &m);
```

passing arguments by *value and return* only
increments m *once*

- Let us consider the code fragment

```
int m = 1;  
doubleincrement(&m, &m);
```

Passing by reference

Caller:

```
...  
int m = 1;  
doubleincrement(&m,&m);  
...  
;
```

m

Pass TWICE the
ADDRESS of m

Variable m gets
incremented
TWICE

Passing by value and result

Caller:

```
...  
int m = 1;  
doubleincrement(&m,&m);  
...  
;
```

m

Pass twice the
VALUE of m:
1 and 1

Return
NEW VALUES:
2 and 2

•
•
•

Passing dynamic types (I)

- Cannot pass dynamic data structures through pointers
- Must send a copy of data structure
- For a linked list
 - Send array with elements of linked list plus unpacking instructions

•
•
•

Passing dynamic types (II)

- We want to pass



- We send to the remote procedure



- Header identifies linked list with four elements

-
-
-

The NYC Cloisters



Rebuilt in NYC from actual cloister stones

•
•
•

Architecture considerations

- The machine representations of floating point numbers and byte ordering conventions can be different:
 - *Little endians* start at the *least* significant byte:
 - *Intel's 80x86 including Pentium*
 - *Big-endians* start at the *most* significant byte:
 - *Sun's SPARC and most RISC processors*

•
•
•

The solution

- Define a network order and convert all numerical variables to that order
 - Use **hton** family of functions
 - *Same as requiring all air traffic control communications to be in English*

•
•
•

Detecting partial failures

- The client must detect *server failures*
 - Can send *are you alive?* messages to the server at fixed time intervals
 - *That is not hard!*

•
•
•

Handling partial executions

- Client must deal with the possibility that the server could have crashed *after* having partially executed the request
 - ATM machine calling the bank computer
 - Was the account debited or not?

•
•
•

First solution (I)

- *Ignore* the problem and *always resubmit* requests that have not been answered
 - Some requests may be executed more than once
- Will work if all requests are *idempotent*
 - Executing them several times has the same effect as executing them exactly once

•
•
•

First solution (II)

- Examples of idempotent requests include:
 - Reading *n* bytes from a fixed location
 - *NOT* reading next *n* bytes
 - Writing *n* bytes starting at a fixed location
 - *NOT* writing *n* bytes starting at current location
- Technique is used by all RPCs in the Sun Microsystems' *Network File System* (NFS)

-
-
-

Second solution

- Attach to each request a *serial number*
 - Server can detect replays of requests it has previously received and refuse to execute them
 - *At most once* semantics
- Cheap but not perfect
 - Some requests could end being partially executed

•
•
•

Third solution

- Use a *transaction mechanism*
 - Guarantees that each request will *either* be *fully executed* or have *no effect*
 - *All or nothing* semantics
- Best and costliest solution
- Use it in all *financial transactions*

-
-
-

An example

- Buying a house using *mortgage money*
 - Cannot get the mortgage without having a title to the house
 - Cannot get title without paying first previous owners
 - Must have the mortgage money to pay them
- Sale is a complex atomic transaction

•
•
•

Realizations (I)

- *Sun RPC:*
 - Developed by Sun Microsystems
 - Used to implement their Network File System
- *MSRPC (Microsoft RPC):*
 - Proprietary version of the DCE/RPC protocol
 - Was used in the Distributed Component Object Model (DCOM).

•
•
•

Realizations (II)

- *SOAP:*
 - Exchanges XML-based messages
 - Runs on the top of HTTP
 - Very portable
 - Very verbose
- *JSON-RPC:*
 - Uses *JavaScript Object Notation* (JSON)